

# Kaggle ML Project using Yelp Dataset

Rohit Gupta

July 26, 2021

```
In [1]: import os
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]: #os.getcwd()
os.chdir("C:\\Users\\Mommy\\Desktop\\Self_Learning\\yelp")
```

```
In [3]: yelp_data = pd.read_csv("yelp.csv")
```

```
In [4]: #yelp_data.head(10)
yelp_data.shape
```

```
Out[4]: (10000, 10)
```

```
In [5]: yelp_data.columns
```

```
Out[5]: Index([u'business_id', u'date', u'review_id', u'stars', u'text', u'type',
              u'user_id', u'cool', u'useful', u'funny'],
              dtype='object')
```

```
In [6]: #yelp_data.head(10)
```

```
In [7]: yelp_data.stars.shape
```

```
Out[7]: (10000L,)
```

```
In [8]: # Filter data frame to have rows with stars being 5 and 1 only. These are the two classes
yelp_data1 = yelp_data[yelp_data['stars']==1]
```

```
In [9]: yelp_data2 = yelp_data[yelp_data['stars']==5]
```

```
In [10]: yelp_data1.shape
```

```
Out[10]: (749, 10)
```

```
In [11]: yelp_data2.shape
```

```
Out[11]: (3337, 10)
```

```

In [12]: yelp_final = pd.merge(yelp_data1, yelp_data2, on='stars')

In [13]: 3337+749

Out[13]: 4086

In [14]: yelp_final = yelp_data.apply(lambda row: row[yelp_data['stars'].isin([5,1])])

In [15]: # Further you can have yelp_final to have only two columns - the review text and the st
# yelp_final.head(10)

In [16]: yelp_final2 = yelp_final[['stars','text']]

In [17]: # Split into training and testing datasets :
# initialize X and Y as series :

X = yelp_final.text
Y = yelp_final.stars

from sklearn.cross_validation import train_test_split

X_train,X_test,Y_train,Y_test = train_test_split(X,Y,random_state=1)

C:\ProgramData\Anaconda2\lib\site-packages\sklearn\cross_validation.py:41: DeprecationWarning: T
"This module will be removed in 0.20.", DeprecationWarning)

In [18]: # check sizes of the train and test datasets
X_train.shape

Out[18]: (3064L,)

In [19]: Y_train.shape

Out[19]: (3064L,)

In [20]: X_test.shape

Out[20]: (1022L,)

In [21]: Y_test.shape

Out[21]: (1022L,)

In [22]: # check if X is series
type(X_train)

Out[22]: pandas.core.series.Series

In [23]: type(X_test)

```

```
Out[23]: pandas.core.series.Series
```

```
In [24]: # import countVectorizer module from sklearn
         from sklearn.feature_extraction.text import CountVectorizer
```

```
         # instantiate the model
         count_vec = CountVectorizer()
```

```
In [25]: # now we need the coun vectorizer model to learn the vocabulary :
         # we need to first learn the vocabulary of the reviews
         # and then create a document term amtrix -
         # with that we will have the mapping of the occurence of each vocab with its occurence
         # go through the row - you build the review/the statement.
```

```
         count_vec.fit(X_train)
```

```
Out[25]: CountVectorizer(analyzer=u'word', binary=False, decode_error=u'strict',
                        dtype=<type 'numpy.int64'>, encoding=u'utf-8', input=u'content',
                        lowercase=True, max_df=1.0, max_features=None, min_df=1,
                        ngram_range=(1, 1), preprocessor=None, stop_words=None,
                        strip_accents=None, token_pattern=u'(?u)\\b\\w\\w+\\b',
                        tokenizer=None, vocabulary=None)
```

```
In [26]: # get feature names or the trained vocab of the review section :
         # count_vec.get_feature_names()
```

```
In [27]: # How to build the doucment term matrix ?
         # use the transform function using . operator on the model you trained.
```

```
         yelp_dtm_train = count_vec.transform(X_train)
         yelp_dtm_test = count_vec.transform(X_test)
         #help(count_vec.transform)
```

```
In [28]: import pandas as pd
         yelp_vect_df_train = pd.DataFrame(yelp_dtm_train.toarray(),columns=count_vec.get_featu
         yelp_vect_df_train.head(3)
```

```
Out[28]:
```

	00	000	00a	00am	00pm	01	02	03	03342	04	...	zucchini	zuchinni	\
0	0	0	0	0	0	0	0	0	0	0	...	0	0	
1	0	0	0	0	0	0	0	0	0	0	...	0	0	
2	0	0	0	0	0	0	0	0	0	0	...	0	0	

	zumba	zupa	zuzu	zwiebel	zzed	éclair	école	ém
0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0	0

```
[3 rows x 16825 columns]
```

```

In [29]: # create the document term matrix for the
yelp_dtm_test = count_vec.transform(X_test)

# create dataframe having column names as trained vocab and rows representing the state
yelp_vect_df_test = pd.DataFrame(yelp_dtm_test.toarray(), columns = count_vec.get_featu

# get shape
yelp_vect_df_test.shape

Out[29]: (1022, 16825)

In [30]: # Now ML begins !!!

In [31]: # now I have the document term matrix - so now using it what to do ?
# well using the document term matrix, we have integer feature representation, so we wi
# along with the star rating.

from sklearn.naive_bayes import MultinomialNB

# instantiate model
mult_nb = MultinomialNB()

# train the data :
mult_nb.fit(yelp_vect_df_train,Y_train)

Out[31]: MultinomialNB(alpha=1.0, class_prior=None, fit_prior=True)

In [32]: # testing the data :
Y_pred_mult_nb = mult_nb.predict(yelp_vect_df_test)

In [33]: print(Y_pred_mult_nb)

[5 5 5 ... 5 1 5]

In [34]: #now we check the accuracy of our model
from sklearn import metrics

metrics.accuracy_score(Y_test,Y_pred_mult_nb)

#0.9187866927592955

Out[34]: 0.9187866927592955

In [35]: # get roc_auc score and get the confusion matrix as well
# Let us print the confusion matrix first :

metrics.confusion_matrix(Y_test,Y_pred_mult_nb)

Out[35]: array([[126,  58],
               [ 25, 813]], dtype=int64)

```

```

In [36]: #help(metrics.classification_report)
         #metrics.classification_report(Y_test,Y_pred_mult_nb)

In [37]: # get the roc_auc score :
         # to get the roc_auc score i need to get the probabilities of the pred class :

         #Y_pred_mult_nb_prob = mult_nb.predict_proba(Y_pred_mult_nb)

In [38]: # Let us use logistic regression
         from sklearn.linear_model import LogisticRegression

         # instantiate the model
         logreg = LogisticRegression()

         # Train the model using the document term matrices
         %time logreg.fit(yelp_vect_df_train,Y_train)

Wall time: 3.71 s

Out[38]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                           intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
                           penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
                           verbose=0, warm_start=False)

In [39]: # test the model :
         Y_pred_logreg = logreg.predict(yelp_vect_df_test)

In [40]: metrics.accuracy_score(Y_test,Y_pred_logreg)

         #0.9256360078277887

Out[40]: 0.9256360078277887

In [41]: # print the confusion matrix, from the accuracy score - logistic regression performs be
         metrics.confusion_matrix(Y_test,Y_pred_logreg)

Out[41]: array([[140,  44],
               [ 32, 806]], dtype=int64)

In [42]: from sklearn import metrics
         metrics.confusion_matrix(Y_test,Y_pred_logreg)

Out[42]: array([[140,  44],
               [ 32, 806]], dtype=int64)

In [43]: #Calculate which 10 tokens are the most predictive of 5-star reviews,
         #and which 10 tokens are the most predictive of 1-star reviews.

         #create a dataframe with token - and the corresponding rating - 5 star or 1 - star.

```

```

#columns : label,tokens ,label_num

# tokens are the feature names
# now for prediction whether the rating is 5 star or 1 star - certain key words will re
# and based on those key words - we get rating. So we can determine occurence of each t
# the 5 star and 1 star rating.

one_star_count = mult_nb.feature_count_[0,:]
five_star_count = mult_nb.feature_count_[1,:]

In [44]: tokens_df2 = pd.DataFrame({'text_review_tokens':yelp_vect_df_train.columns , 'one_star':

In [45]: tokens_df2['one_star'] = tokens_df2.one_star + 1
tokens_df2['five_star'] = tokens_df2.five_star + 1

In [46]: #tokens_df2.head(3)
#five_star_ratio = five_star_count/five_star_class_count [observation/total frequency]
tokens_df2['one_star'] = tokens_df2.one_star/mult_nb.class_count_[0]
tokens_df2['five_star'] = tokens_df2.five_star/mult_nb.class_count_[1]
#tokens_df.sort_values('five_star_ratio',ascending=False).head(3)

# five star to one star ratio :
tokens_df2['five_one_ratio'] = tokens_df2.five_star/tokens_df2.one_star

In [47]: # top ten key words that predict whether the review will be a five star rating.
tokens_df2.sort_values('five_one_ratio',ascending=False).head(3)

Out[47]:
         five_star  one_star text_review_tokens  five_one_ratio
5596      0.077231   0.00354          fantastic      21.817727
10949     0.098039   0.00531             perfect      18.464052
16751     0.024810   0.00177              yum       14.017607

In [48]: # top ten key words that predict whether the review will be a one star rating.
tokens_df2.sort_values('five_one_ratio',ascending=False).tail(3)

Out[48]:
         five_star  one_star text_review_tokens  five_one_ratio
4538      0.0008   0.042478          disgusting      0.018841
12172     0.0004   0.024779             refused      0.016149
14110     0.0004   0.030088          staffperson      0.013299

In [49]: # Visualizations :
# Confusion matrix plot for the multnb model :
conf_matrix_multnb_model = metrics.confusion_matrix(Y_test, Y_pred_mult_nb)

labels = [0, 1]
fig, ax = plt.subplots()
tick_marks = np.arange(len(labels))

```


```

plt.xticks(tick_marks, labels)
plt.yticks(tick_marks, labels)

# create heatmap
sns.heatmap(pd.DataFrame(conf_matrix_multnb_model), annot=True, cmap="YlGnBu", fmt='g')
ax.xaxis.set_label_position("top")
plt.title('Confusion Matrix for Multinomial Naive Bayes Model', y=1.1)
plt.ylabel('True')
plt.xlabel('Predicted')

```

Out[49]: Text(0.5,15, 'Predicted')



In [50]: # Confusion matrix plot for the logistic regression model :

```

conf_matrix_logreg_model = metrics.confusion_matrix(Y_test, Y_pred_logreg)

labels = [0, 1]
fig, ax = plt.subplots()
tick_marks = np.arange(len(labels))
plt.xticks(tick_marks, labels)
plt.yticks(tick_marks, labels)

# create heatmap


```

```

sns.heatmap(pd.DataFrame(conf_matrix_logreg_model), annot=True, cmap="YlGnBu", fmt='g')
ax.xaxis.set_label_position("top")
plt.title('Confusion Matrix for Logistic Regression', y=1.1)
plt.ylabel('True')
plt.xlabel('Predicted')

```

Out[50]: Text(0.5,15, 'Predicted')



```

In [51]: import matplotlib.pyplot as plt
import matplotlib.pyplot as hist
import seaborn as sns

top_ten_fivestar_predictors = tokens_df2.sort_values('five_one_ratio',ascending=False).
top_ten_onestar_predictors = tokens_df2.sort_values('five_one_ratio',ascending=False).t


#plt.plot(top_ten_fivestar_predictors.text_review_tokens[1:10],top_ten_fivestar_predict

sns.set(rc={'figure.figsize':(11.7,8.27)})
sns.barplot(x=top_ten_fivestar_predictors.text_review_tokens[1:10], y=top_ten_fivestar_

```

Out[51]: <matplotlib.axes.\_subplots.AxesSubplot at 0x18d24ba8>






```
In [52]: # Top ten token reviews that predict one star rating :
```

```
sns.set(rc={'figure.figsize':(11.7,8.27)})  
sns.barplot(x=top_ten_onestar_predictors.text_review_tokens[1:10], y=top_ten_onestar_pr
```

```
Out[52]: <matplotlib.axes._subplots.AxesSubplot at 0x1904cda0>
```



```
In [53]: #tokens_df2.text_review_tokens[1:10]
```

```
In [54]: # generate wordcloud for the top ten tokens predicting the 5 stars rating :
```

```
In [55]: # Confusion matrix plot for the logistic regression model :
```

```
In [56]: #5-class classification problem
```

```
#Define X and y using the original DataFrame. (y should contain 5 different classes.)
#Split X and y into training and testing sets.
#Create document-term matrices using CountVectorizer.
#Calculate the testing accuracy of a Multinomial Naive Bayes model.
#Compare the testing accuracy with the null accuracy, and comment on the results.
#Print the confusion matrix, and comment on the results. (This Stack Overflow answer ex
#Print the classification report, and comment on the results.
#If you are unfamiliar with the terminology it uses, research the terms, and then try t
#calculate these metrics manually from the confusion matrix!
```

```
# encountering a problem : my computer cant handle this apparently.
```

```
In [57]: #yelp_final.columns
```

```
In [58]: # initialize the data :
```

```
#X = yelp_data.text
#Y = yelp_data.stars
```

```

In [59]: # initialize training and testing datasets :
        #X_train,X_test,Y_train,Y_test = train_test_split(X,Y,random_state=1)

In [60]: # initialize a new countvectorizer
        #from sklearn.feature_extraction.text import CountVectorizer

        # instantiate new countvec
        #countvec_new = CountVectorizer()

        # train the data
        #countvec_new.fit(X_train)

In [61]: # get the feature names
        #countvec_new.get_feature_names()

        # memory taken alot so the code and output commented.

In [62]: # now convert to get document term matrix for training and testing datasets
        #X_train_dtm_2 = countvec_new.transform(X_train)

In [63]: # look at the array
        #X_train_dtm_2.toarray(5)

In [64]: # get the pandas dataframe having column names as learnt vocab and rows representing th
        #X_train_dtm_df_2 = pd.DataFrame(X_train_dtm_2.toarray(), columns=countvec_new.get_feat
        #

In [65]: #X_train_dtm_df_2.shape

In [66]: # get dtm for test data
        #X_test_dtm_2=countvec_new.transform(X_test)

        # create pandas df :
        #X_test_dtm_df_2 = pd.DataFrame(X_test_dtm_2.toarray(), columns = countvec_new.get_feat

        #X_test_dtm_df_2.head(5)

In [67]: # create the ML model
        #from sklearn.naive_bayes import MultinomialNB

        #mult_nb_2 = MultinomialNB()

        #%time mult_nb_2.fit(X_train_dtm_df_2, Y_train)
        # cannot run 5 class prediction problem - the computer hangs - too resource intensive.

In [68]: #Y_pred_mult_nb_multiclass = mult_nb_2.predict(X_test_dtm_df_2)

```